

Castle Genesis / Rhizoh

External Briefing Pack v1.0

Audience: Policy, governance, operational trust, infrastructure advisors
Tone: Research initiative · infrastructure exploration · controlled deployment experiment
Date: May 2026
Confidentiality: First-contact briefing — not a product pitch or technical repository export.

Document	Contents
Part 1	Executive overview
Part 2	High-level architecture
Part 3	Governance & safety notes
Part 4	Guidance areas

Part 1 — Executive Overview

Castle Genesis is a long-running **infrastructure exploration** program: a structured environment for thinking about how digital worlds, continuity, and AI-assisted interaction can be built **without** treating users as probability fields or hiding how the system behaves.

Rhizoh is the name of the **governance and continuity layer** inside that program — the part that asks: *what is allowed to change state, what may only be observed, and when may the system connect to the outside world?*

This is **not** marketed as “another AI app.” It is closer to a **research-grade control surface**: how to run a world-like experience with explicit boundaries, auditability, and a slow, legal-aware path to any live data connection.

Why did it emerge?

Most platforms optimize for engagement and opaque model upgrades. The motivating question here is different:

Can we build a **non-cheating** digital environment — predictable rhythm, honest limits, and continuity that does not collapse when a vendor model changes?

The work started from art, place, and collaboration (time–place anchors, shared topology) rather than from maximizing session time. Over time it grew into a **serious engineering program**: frozen core execution rules, ingress and legal gates, and formal descriptions of what must *never* leak from “observation” into “control.”

What are we exploring?

Theme	Plain meaning
Continuity	User experience and system state should feel coherent over time, without secret rerolls of “reality.”
Separation of roles	What the user sees, what the system remembers for audit, and what may execute as “truth” are not the same layer.
Controlled external binding	Connecting to real devices or live signals is optional, gated, and off by default until checklists and legal review say otherwise.
Trust without hype	AI is treated as a replaceable motor , not the product. The product is the protocol — what stays stable when models change.

We are **not** claiming a finished global product. We are claiming a **disciplined experiment** with documentation, tests on the control path, and an explicit “ready vs open” decision before any live signal.

Where are we today?

Phase 0.5 — “Safe Reality Layer” (current)

- The **user-facing surface** can be deployed (static hosting, ingress, legal preamble) so the program is **real infrastructure**, not only slides.
- The **data plane** — anything that would ingest live heartbeat or device-like signals from the world — is **intentionally off**. The switch exists in design; it is **not** turned on for production.
- **Perception and copy** have been stabilized so the experience does not feel like a throwaway demo.
- **Formal boundary work** (adversary assumptions, one-way information flow, violation checklists) is written so failures are **named**, not hand-waved.

Next step (only after human “READY”): a controlled Phase 1 probe — at most one minimal signal type, read-only, still forbidden from changing core sealed state.

Why move carefully and involve legal / governance voices?

Risk	Our response
Looking live before it is safe	Deployment ≠ activation; hosting can exist while signals stay off
Observation becoming control	Witness logs may grow; core state must not follow
Regulatory surprise (EU-facing users)	Legal preamble, counsel-oriented packs, minimal Phase 1 signal design
“AI governance” as marketing	Moderation and safety as operational duties


One sentence summary

Castle Genesis / Rhizoh is a carefully bounded infrastructure experiment: a world-like surface with a frozen control ethic, no live data-plane activation until explicit readiness, and a governance posture that prefers slow trust to fast spectacle.

Part 2 — High-Level Architecture

Purpose: This is a **thought-through architecture**, not only a narrative. Blocks only — no API lists or repository map.

Layer flow (diagram summary)



Layer roles (30 seconds)

Layer	Role in one line
Interaction	What people see and enter through (including legal boundary screens).
Orchestration	How modes, routes, and control decisions are sequenced — not free-form scripting.
AI	Models assist and narrate; they do not silently become the authority for sealed state.
Memory / continuity	Core truth, observation logs, and display are kept apart.

Layer	Role in one line
Realtime	Live transport where the experiment requires it — under orchestration and policy.
Governance / safety	Law, activation switches, moderation mindset, operational responsibility.

Critical design choice (today): flow from **observation** back into **core** is **forbidden**. External signals, when eventually allowed, are designed to **verify presence**, not to rewrite sealed ledger state.

Current phase: Surface and control path — active and testable. Live world signal ingestion — **off** until explicit readiness and legal alignment.

Part 3 — Governance & Safety Notes

Controlled rollout

Decision	Meaning
Ready?	Infrastructure, legal pack, and automated checks say we <i>may</i> operate responsibly.
Open?	Any live external signal is allowed — forbidden until Ready is explicitly signed.

Hosting a static site is **not** the same as activating data collection. Rollout is **phased**: surface first, minimal read-only signal later.

Security awareness

- Least privilege; credentials out of scope for first briefings.
- Data-plane ingestion **designed but not activated** for production.
- Adversary scenarios with **named violation types**.
- Interpretation constraints: external input must not find a side door into sealed state.

Moderation and harm mindset

- Harmful use anticipated at product and community level.
- Moderation as **ongoing operational duty**, not a one-time safety badge.
- AI surface copy is **non-authoritative** relative to sealed state and legal text.

Data minimization (Phase 1 design intent)

- Opaque device/session references — no names in pipeline design.
- Liveness, build id, consent epoch — no free-text payload.
- Rate limits and schema validation at gateway.
- No GPS or rich telemetry in default Phase 1 class without separate legal approval.
- Observation logs may grow; **core sealed state must not** grow from those logs.

User trust approach

- Honest baseline: avoid variable-reward “reality casinos.”
- Ingress transparency: terms and privacy where required.
- No fake social proof in production UI.
- Users should understand the system **observes under rules**, not that it secretly runs them.

Operational responsibility

- Checklists with human **READY / HOLD** sign-off.

- Exportable ops run artifacts (readiness, violation harnesses).
- Culture documents separate from executable freeze.

Closing: We ask whether phased activation, data minimization, non-interference rules, and operational checklists match what you expect before broader exposure.

Part 4 — Guidance Areas

We are not requesting open-ended consulting. We identify **specific lenses** where informed perspective would reduce risk. Partial guidance is welcome.

We are currently looking for guidance regarding:

Publication & communication readiness — research / controlled experiment framing; what not to claim externally; how to describe AI without implying models are the system of record.

Legal exposure (EU-facing awareness) — high-level KVKK / GDPR-aligned posture for ingress and minimal future heartbeat class; whether “deployment without activation” is useful for counsel; timing of formal counsel vs READY sign-off.

AI governance considerations — human oversight, documentation, change control; non-binding AI output vs sealed core state; red flags on vendor/model churn.

EU-facing regulatory awareness — whether phased, opt-in, data-minimized signal direction is directionally sane; categories to exclude in Phase 1; reference frameworks for documenting decisions (not compliance certificates).

Infrastructure trust expectations — evidence before linking identity surfaces; credibility of two-gate (ready vs open) model; staging vs production communication norms.

Cybersecurity considerations — usefulness of violation taxonomy for security reviewers; responsible disclosure posture; gaps for static+ingress without live ingest.

Organizational structure for future scaling — who signs READY/HOLD; legal/ops/engineering separation; advisory board or counsel retainer before scale; avoiding implicit DPO-like duties on founders.

Helpful output formats

Format	Useful?
Short written reactions (email / bullet memo)	✓ Preferred
30–60 min conversation after Part 1	✓
Full repository audit	✗ Not in first contact
Compliance certification	✗ Out of scope for this phase

Not included in first contact

API secrets, internal prompts, full manifesto archive, agent internals, production endpoint details, orchestration internals.

